

Regular Expression

COMP3220 – Principle of Programming Languages

Zhitao Gong

2016 Spring

Outline

Combining Languages

Regular Language

Regular Expression in Theory

Regular Expression in Practice

Summary

Concatenation of Languages

The concatenation of two languages L_1 and L_2 is

$$L_1 L_2 = \{xy \mid x \in L_1 \text{ and } y \in L_2\}$$

Basically, concatenate every string from the first language with every string from the second. E.g., if $L_1 = \{a, b\}$ and $L_2 = \{c, d\}$, then $L_1 L_2 = \{ac, ad, bc, bd\}$.

Union of Languages

The union of two languages L_1 and L_2 is

$$L_1 \cup L_2 = \{x \mid x \in L_1 \text{ or } x \in L_2\}$$

Basically, it is a plain *set union*.

Kleene Closure of A Language

Given a language L , its Kleene closure L^* may be defined recursively.

Base clause $\epsilon \in L^*$

Inductive clause $xw \in L^* \forall x \in L^* \text{ and } w \in L$

Extremal clause A string is in L if and only if it can be obtained from the above two clauses.

E.g., if $L = \{ab, cd\}$,

$L^* = \{\epsilon, ab, cd, abab, abcd, cdab, cdcd, ababab, \dots\}$.

Outline

Combining Languages

Regular Language

Regular Expression in Theory

Regular Expression in Practice

Summary

Recursive Definition

The regular language over an alphabet Σ may be defined *recursively*.

Basic clause \emptyset , $\{\epsilon\}$, $\{a\} \forall a \in \Sigma$ are regular language.

Inductive clause If L_r and L_s are regular languages, then $L_r \cup L_s$, $L_r L_s$, L_r^* are regular languages.

Extremal clause A language is regular if and only if it can be obtained from the above two clauses.

Example

Given an alphabet $\Sigma = \{a, b\}$

- ▶ $\{a\}$ and $\{b\}$ are regular language.
- ▶ $\{a, b\}$ ($= \{a\} \cup \{b\}$), $\{ab\}$ ($= \{a\}\{b\}$) are regular language.
- ▶ $\{a\}^*$ ($= \{\epsilon, a, aa, aaa, \dots\}$) is regular language.
- ▶ Σ^* is also a regular language since $\{a, b\}$ is regular.

Outline

Combining Languages

Regular Language

Regular Expression in Theory

Regular Expression in Practice

Summary

Recursive Definition

Theorem

A regular expression is a string r that denotes a regular language $L(r)$ over some alphabet Σ .

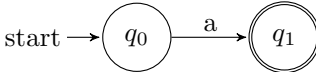
Basic clause \emptyset , ϵ and a are regular expressions denoting \emptyset , $\{\epsilon\}$ and $\{a\} \forall a \in \Sigma$, respectively. They are called *atomic* regular expression.

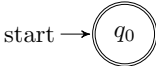
Inductive clause r and s are regular expressions denoting regular language L_r and L_s , then the follows are *compound* regular expressions

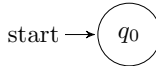
- ▶ $r + s$ denotes $L_r \cup L_s$
- ▶ rs denotes $L_r L_s$
- ▶ r^* denotes L_r^*

Extremal clause A string is a regular expression if and only if it can be obtained from above two clauses.

Base Clause NFAs

- ▶ $\forall a \in \Sigma, \mathcal{L}(a) = a,$ 

```
graph LR; start((start)) --> q0((q0)); q0 -- a --> q1(((q1)))
```
- ▶ $\mathcal{L}(\epsilon) = \epsilon,$ 

```
graph LR; start((start)) --> q0(((q0)))
```
- ▶ $\mathcal{L}(\emptyset) = \emptyset,$ 

```
graph LR; start((start)) --> q0((q0))
```

Inductive Clause – Union

M_0, M_1 recognize $\mathcal{L}(R_0)$ and $\mathcal{L}(R_1)$ respectively.




Figure: $R = R_0 + R_1$

Inductive Clause – Concatenation

M_0 , M_1 recognize $\mathcal{L}(R_0)$ and $\mathcal{L}(R_1)$ respectively.




Figure: $R = R_0R_1$

Inductive Clause – Kleene Closure

M_0 recognizes $\mathcal{L}(R_0)$




Figure: $R = R_0^*$

Equivalence of RE and FSA

1. We need to show for every RE, there is a FSA that accepts the language. Done with NFA- ϵ .
2. We need to show for every FSA, there is a RE denoting its language. Trickier and ignored for now.

Regular Expression Conventions

1. For precedence, $(R) > R^* > R_1 R_2 > R_1 + R_2$.
 - ▶ $(a + (b(c^*)))$ is the same as $a + bc^*$.
 - ▶ $ab^*c + d$ is the same as $((a(b^*))c) + d$.
2. Similar to string concatenation, concatenation of k regular expression r 's is written as $r^k = \overbrace{rr \dots r}^k$, specifically $r^0 = \emptyset$.

Regular Expression Examples

- ▶ $\mathcal{L}(01) = \{01\}$
- ▶ $\mathcal{L}(01 + 0) = \{01, 0\}$
- ▶ $\mathcal{L}(0(1 + 0)) = \{01, 00\}$
- ▶ $\mathcal{L}(0^*) = \{\epsilon, 0, 00, 000, \dots\}$
- ▶ $\mathcal{L}((0 + 01)^*(\epsilon + 1)) =$

Regular Expression Examples

- ▶ $\mathcal{L}(01) = \{01\}$
- ▶ $\mathcal{L}(01 + 0) = \{01, 0\}$
- ▶ $\mathcal{L}(0(1 + 0)) = \{01, 00\}$
- ▶ $\mathcal{L}(0^*) = \{\epsilon, 0, 00, 000, \dots\}$
- ▶ $\mathcal{L}((0 + 01)^*(\epsilon + 1))$ = all strings of 0's and 1's without two consecutive 1's.

Why It Matters

Theorem

For any regular expression R , there is FSA M such that

$$\mathcal{L}(R) = \mathcal{L}(M)$$

Why It Matters

Theorem

For any regular expression R , there is FSA M such that

$$\mathcal{L}(R) = \mathcal{L}(M)$$

Working solution for matching regular expressions against text.

- ▶ Convert regular expression to an NFA.
- ▶ (Optionally) convert NFA to a DFA.
- ▶ Run the text through the DFA.

Example Regular Expression \rightarrow NFA




Figure: $L = ab \cup a^*$

- ▶ a : $\text{start} \rightarrow \text{circle} \xrightarrow{a} \text{final circle}$, b : $\text{start} \rightarrow \text{circle} \xrightarrow{b} \text{final circle}$
- ▶ ab : $\text{start} \rightarrow \text{circle} \xrightarrow{a} \text{circle} \xrightarrow{\epsilon} \text{circle} \xrightarrow{b} \text{final circle}$
- ▶ a^* : $\text{start} \rightarrow \text{final circle} \xrightarrow{\epsilon} \text{circle} \xrightarrow{a} \text{final circle} \xrightarrow{\epsilon} \text{circle}$

D/N-FA/Regular Expression Accept Regular Languages

\mathcal{L}_{DFA} All languages accepted by DFAs

\mathcal{L}_{NFA} All languages accepted by NFAs

\mathcal{L}_{REG} All languages denoted by regular expressions

\mathcal{L}_R Regular Languages

$$\mathcal{L}_{DFA} = \mathcal{L}_{NFA} = \mathcal{L}_{REG} = \mathcal{L}_R$$

Outline

Combining Languages

Regular Language

Regular Expression in Theory

Regular Expression in Practice

Summary

Regular Expression Caveat

Some people, when confronted with a problem, think "I know I'll use regular expressions." Now they have two problems.

– Jamie Zawinski (flame war on alt.religion.emacs)

Writing RE is like writing a program.

- ▶ Need to understand programming model.
- ▶ Can be easier to write than read.
- ▶ Can be difficult to debug.

Are You Smart Enough

<http://www.ex-parrot.com/~pdw/Mail-RFC822-Address.html>

[illegible]

RE Standard

IEEE POSIX standard has three set of compliance:

- ▶ *B*-asic RE (BRE)
- ▶ *E*-xtended RE (ERE) and
- ▶ ~~*S*-imple RE (SRE)~~ (deprecated).

ERE extends BRE

1. it adds `?`, `+`, `|`, and
2. removes the need to escape *meta-characters* `()` and `{}`.

Notations

- ▶ Regular expression is abbreviated as *regex* or *RE* in the following discussions.
- ▶ We use `regex` to denote a regex, `text` a plain text string.
- ▶ `regex` = `text` denotes that `regex` matches `text`.
- ▶ `regex` \neq `text` denotes that `regex` does not match `text`.
- ▶ `regex1` = `regex2` denotes that both regexes are equivalent.

Literal Characters

Single literal character, such as `a`, matches the *first occurrence* of the character.

- ▶ Given the string `Jack is a boy` it matches `a` after the `J`.
- ▶ The fact that `a` is in middle of the string does not matter in this case unless you use *word delimiter*.

Matching is *case sensitive* by default. E.g., `cat` \neq `Cat`. It is configurable however.

Meta Characters

- ▶ escape, backslash \
- ▶ anchors, caret ^, dollar sign \$
- ▶ matches all (almost), period .
- ▶ alternation, vertical bar |
- ▶ question mark ?, asterisk *, plus sign +
- ▶ grouping, opening and closing parenthesis (,)
- ▶ character set, opening square bracket [,] and
- ▶ repetition, opening curly brace {, }.

Meta Character Match

To match meta characters, you need to *escape* them with a `\`.

- ▶ to match `1+1=2`, the correct regex is `1\+1=2`.
- ▶ `1+1=2` is a valid regex, matching `111=2`, `1111=2`, `123+111=234`, etc.

Non-Printable Characters

Conventional formatter, `\t` (tab, 0x09), `\r` (carriage return, 0x0D), `\n` (line feed, 0x0A) and etc. The followings are inconsistent.

- ▶ Control sequences, `\cA` = Control+A.
- ▶ Unicode characters, `\u20AC` (or `\x{20AC}`) = €.

Character Class

With a "character class", also called "character set", you can tell the regex engine to match *only one* out of several characters.

- ▶ `[ae]` = `a` or `e`. E.g., `gr[ae]y` = `gray` or `grey`. The order in the bracket *does not* matter.
- ▶ Hyphen specifies a range.
 - ▶ `[0-9]` = *single* digit between 0 and 9
 - ▶ `[a-z]` = *single* character between a and z.
 - ▶ `0[xX][A-Fa-f0-9]+` = C-style hexadecimal number.

Negated Character Class

A caret `^` after the opening square bracket *negates* the character class. The result is that the character class matches any character that is *not* in the character class.

- ▶ `[^0-9\r\n]` = any single character that is *not* a digit or a line break.
- ▶ `q[^u]` does *not* mean "a `q` not followed by a `u`". Instead it means "a `q` followed by a character that is not a `u`". E.g.,
`q[^u] ≠ q`, `q[^u] = q`, note that space after `q`.

Meta Character inside Character Class

Inside character class, only `]`, `^`, `\` and `-` are special.

- ▶ `[\]` = `]`, `[\^]` = `^`, etc.
- ▶ `[\+]` = `[+]` = `+`.
- ▶ `[0^]` = `0` or `^`
- ▶ The hyphen can be included a) right after the opening bracket, or b) right before the closing bracket, or c) right after the negating caret. E.g., `[-x]` = `[x-]` = `x` or `-`.

Shorthand Character Class

Some character classes are often used.

- ▶ digit `\d` = `[0-9]`
- ▶ word character `\w` = `[A-Za-z0-9_]`
- ▶ whitespace character `\s` = `[\t\r\n\f]`.
- ▶ Shorthand may also be used inside square brackets. E.g., `[\s\d]` = a whitespace or a digit.

Negated Shorthand Character Class

Most shorthand character classes have negated shorthand.

- ▶ non-digit `\D` = `[^\d]`
- ▶ non-whitespace `\S` = `[^\s]`
- ▶ non-word character `\W` = `[^\w]`
- ▶ `[\S\D]` \neq `[^\s\d]`

Dot

Dot `.` matches any character *but line breaks*.

Given the four date formats

- ▶ `mm/dd/yy`
- ▶ `mm-dd-yy`
- ▶ `mm.dd.yy`
- ▶ `mm dd yy`
- ▶ A quick solution is `\d\d.\d\d.\d\d`.

Dot

Dot `.` matches any character *but line breaks*.

Given the four date formats

- ▶ `mm/dd/yy`
- ▶ `mm-dd-yy`
- ▶ `mm.dd.yy`
- ▶ `mm dd yy`
- ▶ A quick solution is `\d\d.\d\d.\d\d`. However `02512703` is also a match.

Dot

Dot `.` matches any character *but line breaks*.

Given the four date formats

- ▶ `mm/dd/yy`
- ▶ `mm-dd-yy`
- ▶ `mm.dd.yy`
- ▶ `mm dd yy`
- ▶ A quick solution is `\d\d.\d\d.\d\d`. However `02512703` is also a match.
- ▶ Modified version is `\d\d[- /.]\d\d[- /.]\d\d`.

Dot

Dot `.` matches any character *but line breaks*.

Given the four date formats

- ▶ `mm/dd/yy`
- ▶ `mm-dd-yy`
- ▶ `mm.dd.yy`
- ▶ `mm dd yy`
- ▶ A quick solution is `\d\d.\d\d.\d\d`. However `02512703` is also a match.
- ▶ Modified version is `\d\d[- /.]\d\d[- /.]\d\d`. However it matches `99/99/99`.

Anchors

Anchors match positions instead of characters.

- ▶ `^` = the position before the first character in the string. E.g.,
`^a` = `a` in `abc`, `^b` \neq `b` in `abc`.
- ▶ `$` = the position right after the last character in the string.
E.g., `c$` = `c` in `abc`, `a$` \neq `a` in `abc`.
- ▶ In *multi-line mode*, `^` and `$` matches the start and end of each line respectively.

Word Boundary

`\b` matches a position called "word boundary". It allows whole-words-only search. Its negation is `\B`.

- ▶ Before the first character in the string, if the first character is a word character.
- ▶ After the last character in the string, if the last character is a word character.
- ▶ Between two characters in the string, where one is a word character and the other is not a word character.

E.g., `\b4\b` \neq `4` in `44 sheets of a4`.

Alternation

Alternation `|` is used to match a single regular expression out of several possible ones.

- ▶ `cat|dog` = `cat` or `dog`,
- ▶ `cat|dog|mouse|fish` = `cat` or `dog` or `mouse` or `fish`

The `|` has the *lowest precedence* of all regex operators.

- ▶ `\bcat|dog\b` = `cat` in `Hi cats!!` or `dog` in `hotdog!`.
- ▶ `\b(cat|dog)\b` matches nothing in `Hi cats!!` or `hotdog!`.

Optional Item

The question mark ? the preceding token optional. The question mark is called a *quantifier*.

- ▶ `colou?r` = `color` and `colour`.
- ▶ `Nov(ember)?` = `Nov` and `November`.
- ▶ `Feb 23(rd)?` = `Feb 23rd` in `Today is Feb 23rd` since `?` is *greedy*, discussed later.

Repetition

Two more quantifiers.

- ▶ * tells the regex engine to attempt to match the preceding token zero or more times.
- ▶ + tells the regex engine to attempt to match the preceding token once or more.

E.g., `<[A-Za-z][A-Za-z0-9]*>` = an HTML tag without attributes, i.e., ``, `<html>`, etc.

Limiting Repetition

An additional quantifier allows you to specify how many times a token can be repeated, the syntax is `{min, max}`.

▶ `0?` = `0{0,1}`

▶ `0+` = `0{1,}`

▶ `0*` = `0{0,}`

▶ `0{3}` = `000`

Greedy and Lazy Quantifiers

?, *, +, {min,max} are all greedy, they tell the regex engine to match as many as possible unless it results in a match fail.

- ▶ `<.+>` = `emph` in a `emph` test, not `` or ``.

Putting a ? after these greedy quantifiers makes them *lazy* (or *ungreedy*, *reluctant*).

- ▶ `<.+?>` = `` and `` in a `emph` test.

Atomic Grouping

An **atomic group** is a group that, when the regex engine exits from it, automatically throws away all backtracking positions remembered by any tokens inside the group.

- ▶ Atomic groups are non-capturing.
- ▶ The syntax is `(?>group)`

- ▶ `a(bc|b)c` = `abcc` and `abc`.
- ▶ `a(?>bc|b)c` = `abcc`
- ▶ `a(?>bc|b)c` \neq `abc`
- ▶ `\b(integer|insert|in)\b` \neq `integers`. Optimized version is `\b(?>integer|insert|in)\b`

Possessive Quantifiers

Possessive quantifiers are a syntax sugar to place an atomic group around a single quantifier. `X*+` = `(?>X*)`.

- ▶ `".*"` = `"abc"` in `"abc"x`
- ▶ `".*+"` \neq `"abc"` in `"abc"x`

Grouping and Capturing

- ▶ Only parentheses, (and) can be used for grouping, which allows you to apply a quantifier to the entire group or to restrict alternation to part of the regex.
 - ▶ `cat|dog` = `cat` and `dog` in `rains cats and dogs`.
 - ▶ `ca(t|d)og` = `catog` or `cadog`.
- ▶ () not only creates groups, but *numbered capturing groups*.
 - ▶ `get(value)\1` = `getvaluevalue`.
- ▶ To disable capturing, put `?:` right after opening (.
 - ▶ `get(?:value)\1` results in error.

Backreference

Previously, `get(value)\1` = `getvaluevalue`.

- ▶ `\1` is called backreference, matching the *exact same text* that was matched by the *first capturing group*.
- ▶ `\2` matches the *second capturing group*, etc.
- ▶ Backreference may be reused many times.

Backreference stores the last match.

- ▶ `([abc]+)=\1` = `cab=cab`
- ▶ `([abc])+=\1` \neq `cab=cab`
- ▶ `([abc])+=\1` = `cab=b`

`\b(\w+)\s+\1\b` = doubled word, e.g. `the the`.

Backreference to Failed Groups

- ▶ $(q?)b\backslash 1 = b$ $q?$ matches nothing, so does $\backslash 1$.
- ▶ $(q)?b\backslash 1 \neq b$ $(q)?$ proceeds to b , then $\backslash 1$ which fails by mimicking the result of the group.
- ▶ $(q)?b\backslash 1? = b$

Regex Engine is *Eager*

Regex engine stops searching as soon as it finds a valid match.

- ▶ `Get|GetValue` = `Get` in `GetValue`.
- ▶ `GetValue|Get` = `GetValue` in `GetValue`.
- ▶ `\b(Get|GetValue)\b` = `GetValue` in `GetValue`
- ▶ `Get(Value)?` = `GetValue` in `GetValue` since `?` is *greedy*.

Advanced Topics

- ▶ Lookahead and lookahead
- ▶ Conditionals
- ▶ Recursion
- ▶ Subroutine
- ▶ etc.

Outline

Combining Languages

Regular Language

Regular Expression in Theory

Regular Expression in Practice

Summary

Regular Expression and Regular Language

1. RE is used to *denote* regular language.
2. RE may be implemented using NFA.
3. Without knowing regular language, RE is also useful
4. RE is crazy.